

PRIE: Placement Readiness Intelligence Engine with Calibrated Predictive Modeling and Constrained Prescriptive Recourse

Sivasubramanian R
Department of AIML
Malla Reddy University
Dulapally, Hyderabad, India
Sivasubramanian243@gmail.com

Kundala Dhana Naga Shankar
Department of AIML
Malla Reddy University
Dulapally, Hyderabad, India
Kundaladhana2004@gmail.com

Kavala Rajeev
Department of AIML
Malla Reddy University
Dulapally, Hyderabad, India
Kavalaraajeev@gmail.com

Kouru Rudra Teja
Department of AIML
Malla Reddy University
Dulapally, Hyderabad, India
Rudrateja08@gmail.com

Abstract—The transition from undergraduate engineering education to professional technical employment is severely hindered by the fragmentation of campus placement preparation systems. Conventional educational data mining approaches predominantly deploy retrospective, uncalibrated classifiers that predict point-in-time placement outcomes without providing actionable pedagogical recourse. In this paper, we propose the Placement Readiness Intelligence Engine (PRIE), an integrated continuous intelligence framework that unifies heterogeneous multi-source student telemetry into a standardized 22-dimensional Student Profile Vector ($x_{spv} \in [0.0, 1.0]^{22}$) paired with an explicit observation mask. PRIE deploys a cost-sensitive gradient boosted decision tree classifier coupled with Platt probability scaling to produce well-calibrated placement readiness probabilities. To overcome algorithmic opacity, polynomial-time TreeSHAP isolates diagnostic feature contributions, while constrained Diverse Counterfactual Explanations (DiCE) generate sparse, actionable recourse paths ($k = 2.47 \pm 0.52 \leq 3.0$ mutable features modified) that maintain 100.0% invariance across immutable institutional department attributes (F_{17}). Identified competency deficits trigger Kahn’s topological scheduler over a 38-node computer science concept directed acyclic graph (DAG), achieving a 0.0% prerequisite precedence violation rate across generated curricula. Furthermore, weighted tri-modal late fusion across acoustic prosody, video composure, and speech clarity dampens single-sensor diagnostic variance by $77.98\% \pm 3.99\%$ ($t = 9.88, p = 0.0022$) while operating within a 1,120 ms conversational turn-taking latency budget. Across a 5-seed evaluation on synthetic engineering cohorts ($N = 2,500$), the proposed calibrated model achieves a test accuracy of $95.20\% \pm 1.17\%$ (94.60% in hold-out test partition), an ROC-AUC of 0.9922 ± 0.0038 , an Expected Calibration Error (ECE) of 0.0350 ± 0.0057 , and a Brier score of 0.0339 ± 0.0096 .

Index Terms—Placement readiness, Educational data mining, Explainable artificial intelligence, Student Profile Vector, Platt calibration, Algorithmic recourse, Multimodal fusion, Kahn topological sort.

I. Introduction

The transition of engineering undergraduates into professional technical careers represents a foundational milestone for student socio-economic mobility and institutional workforce alignment [1]. However, higher education institutions globally confront an acute readiness crisis: while hundreds of thousands of engineering students participate in campus recruitment drives annually, technology employers consistently report severe competency deficits in practical software design, clean algorithmic problem-solving, architectural debugging, and professional workplace communication [2]. In conventional university placement preparation, institutional triage relies almost exclusively on static administrative metrics—principally cumulative Grade Point Average (CGPA) or terminal semester examination marks [3].

Nevertheless, static academic grades represent lagging indicators that correlate weakly with modern agile industry requirements [4]. Furthermore, campus placement preparation remains fragmented into disconnected software silos: across our systematic review of the 44 verified career-readiness and educational analytics systems analyzed in our research foundation, 42 systems (95.5%, 42/44) address only one or two dimensions in isolation without an integrated continuous student state [6]. Students typically interact with standalone Applicant Tracking System (ATS) resume checkers that perform rudimentary keyword counting, separate coding portals that evaluate unit test pass rates without assessing architectural design quality, and uncalibrated conversational bots that dispense generic interview advice [3], [5].

From a machine learning perspective, existing educational data mining (EDM) models exhibit five debilitating

limitations:

- 1) Retrospective Point-in-Time Evaluation: Models predominantly operate as post-hoc binary classifiers in final semesters, precluding early multi-semester diagnostic intervention [6], [7].
- 2) Probabilistic Miscalibration: Classifiers output over-confident risk scores that fail to reflect empirical posterior probabilities, distorting high-stakes advising triage [8].
- 3) The Descriptive-to-Prescriptive Divide: Explainable AI (XAI) frameworks deploy global feature importances or post-hoc Shapley attributions that describe historical deficits but offer zero computable recourse over actionable, mutable variables [16], [17].
- 4) Sensory Volatility in Mock Interviews: Single-sensor conversational assessment tools suffer severe diagnostic variance and high conversational turn latency [12], [13], [19].
- 5) Prerequisite Blindness: Recommendation algorithms suggest unsequenced study topics that violate fundamental curricular prerequisite dependencies [14], [20].

To overcome these challenges, this paper presents the Placement Readiness Intelligence Engine (PRIE), an integrated continuous intelligence architecture designed for multimodal student tracking, calibrated probability estimation, and closed-loop prescriptive remediation. The research is guided by six foundational research questions (RQs):

- RQ1: Can spatial document representations overcome visual layout destruction in multi-column resume parsing?
- RQ2: Can tri-modal late fusion reduce diagnostic volatility in automated mock interviews within a sub-1.5-second conversational latency budget?
- RQ3: Does probability calibration produce reliable, trustworthy placement risk scores compared to uncalibrated baselines?
- RQ4: Can constrained counterfactual optimization generate sparse, actionable remediation paths while strictly preserving immutable demographic attributes?
- RQ5: How effectively does topological graph scheduling enforce prerequisite precedence constraints in personalized curricula?
- RQ6: Can cosine similarity gating reliably reject out-of-domain queries and prompt injections prior to curriculum retrieval context augmentation?

The primary scientific and engineering contributions of this work are summarized as follows:

- 1) Unified Latent State Representation: We formulate the 22-dimensional Student Profile Vector ($x_{\text{spv}} \in [0.0, 1.0]^{22}$) paired with an explicit observation mask ($m \in \{0, 1\}^{22}$) to harmonize academic transcripts,

practical coding scores, 2D resume metrics, and paralinguistic interview telemetry.

- 2) Calibrated Predictive Modeling: We demonstrate that integrating cost-sensitive gradient boosted decision trees with Platt probability scaling contracts Expected Calibration Error (ECE) to 0.0350 ± 0.0057 and Brier score loss to 0.0339 ± 0.0096 , reducing probability calibration error under the evaluated synthetic-cohort conditions.
- 3) Constrained Prescriptive Recourse: We bridge the descriptive-to-prescriptive divide by coupling polynomial-time TreeSHAP ($O(TLD^2)$) with constrained Diverse Counterfactual Explanations (DiCE), generating sparse recourse paths ($k = 2.47 \pm 0.52 \leq 3.0$) that maintain 100.0% invariance across immutable institutional attributes across all evaluated profiles.
- 4) Multimodal Variance Damping & Topological Remediation: We observe that weighted tri-modal late fusion dampens interview scoring variance by $77.98\% \pm 3.99\%$ ($t = 9.88, p = 0.0022$) under a 1,120 ms turn latency budget, while Kahn’s topological scheduler enforces prerequisite precedence constraints (0.0% violation rate) over a 38-node computer science concept DAG.

The remainder of this paper is organized as follows: Section II presents a thematic synthesis of related work and establishes the formal research gap. Section III details the PRIE system architecture. Section IV formalizes the scientific methodology. Section V outlines the experimental design. Section VI presents empirical results across all research questions. Section VII discusses theoretical implications and trade-offs. Section VIII examines limitations and threats to validity. Section IX concludes the paper and outlines future directions.

II. Related Work and Research Gap

Graduate employability prediction, educational data mining, and multimodal career assessment have received substantial scholarly attention. Rather than enumerating studies sequentially, this section synthesizes the literature thematically across six core dimensions.

A. Thematic Literature Synthesis

- 1) Employability and Student Performance Prediction: Early EDM investigations formulated employability as supervised binary or multinomial classification over terminal graduation records. Casuat and Festijo [7] applied decision trees and ensemble techniques to institutional records, achieving 84.50% classification accuracy. Rao and Swamy [8] benchmarked classical classifiers, reporting 78.40% accuracy on academic records. Olipas [1] investigated machine learning models for career readiness, achieving 88.40% accuracy using Random Forests. Patel and Nair [4] incorporated psychometric indicators with academic scores, attaining 91.20% accuracy. While these studies

achieve competitive nominal accuracy, they exhibit severe probabilistic miscalibration: raw confidence scores do not reflect empirical placement probabilities, compromising institutional advising triage.

2) Educational Learning Analytics & Temporal Telemetry: Van Wyk and Du Plessis [2] investigated AI-driven learning analytics in higher education, highlighting that continuous telemetry tracking (e.g., login cadence, formative test attempts, and platform engagement intensity) yields substantially stronger predictive utility than point-in-time administrative marks. Chen and Hwang [5] conducted a bibliometric analysis of artificial intelligence in education, emphasizing that institutional adoption requires algorithmic explainability, fairness guarantees, and transparent probability calibration rather than black-box margin scores.

3) Explainable AI and Algorithmic Recourse: Recent works have introduced post-hoc interpretability into academic early-warning systems. Hidayatulloh et al. [16] applied Shapley Additive Explanations (SHAP) to student performance models, showing that feature attribution enhances educator trust. Similarly, Joshi and Kulkarni [17] deployed LightGBM and TreeSHAP to isolate risk factors across engineering branches. Nevertheless, these frameworks remain confined to descriptive explanation: they explain why a student is predicted to fail (e.g., attributing risk to a low historical GPA or high school grade) but provide no prescriptive recourse, because completed academic records cannot be modified retroactively [16], [17]. Prescriptive intervention requires solving constrained optimization problems over actionable, mutable variables.

4) Automated Resume Parsing & Information Extraction: Traditional resume screening relies on flat text tokenizers or keyword matching [9], [10]. Verma and Mehta [15] developed ResuMatch, which applied semantic sentence embeddings to compare resume text with job descriptions. Zhang et al. [11] introduced Career-gAide, emphasizing career path recommendation based on resume profiles. However, standard extraction pipelines discard 2D visual layout coordinates, causing column interleaving errors in modern multi-column CVs.

5) Automated Mock Interview Assessment: AI-driven interview assessment has progressed from single-channel acoustic scoring to multimodal architectures. Deshmukh and Kulkarni [12] reviewed automated interview systems, noting that single-sensor evaluations suffer from acoustic and visual noise. The Advanced Innovation Consortium [13] demonstrated that fusing facial expression recognition, speech emotion analysis, and textual NLP produces more holistic candidate assessments. Srinivasan and Radhakrishnan [19] proposed a voice-driven interview simulator pairing Whisper speech recognition with LLM questioning, but observed conversational turn latency exceeding 2.5 seconds.

6) Knowledge Graphs, RAG & Curricular Scheduling: Tan et al. [14] addressed personalized learning path-

ways, demonstrating that automated curriculum recommendations must adhere to course prerequisite DAGs to avoid cognitive overload. Sutherland and Miller [18] surveyed retrieval-augmented generation (RAG) in education, demonstrating that similarity threshold gating rejects out-of-domain queries prior to context augmentation. Fernandez and Gomez [20] reviewed automated question generation, showing that causal concept graph grounding is essential for generating psychometrically valid diagnostic items. Finally, Babureddy and Mathew [21] proposed a triangular digital twin concept linking student, faculty, and industry requirements.

B. Research Gap Formulation

Cross-analysis of the existing literature reveals three structural voids:

- 1) The Single-Module Isolation Chasm: Across our systematic review of the 44 verified career-readiness systems analyzed in our research foundation, 42 systems (95.5%, 42/44) evaluate only one or two functional dimensions in isolation (e.g., ATS parsing alone or terminal placement classification alone) without an integrated continuous latent state [3], [6].
- 2) The Descriptive-to-Prescriptive Divide: Existing XAI models output historical attributions but fail to formulate constrained mathematical recourse paths bounded by student cognitive budgets [16], [17].
- 3) Probabilistic Miscalibration: Classifiers optimize unweighted loss functions, producing overconfident probabilities that undermine algorithmic credibility in high-stakes educational advising [8].

PRIE directly resolves these gaps by uniting multimodal telemetry into an invariant 22-dimensional Student Profile Vector (x_{spv}), coupling calibrated gradient boosting with constrained counterfactual recourse, and enforcing topological graph sequencing over personalized remediation milestones.

III. PRIE System Architecture

Grounded in the Phase 05 system architecture specifications, PRIE is organized as a modular, four-tier continuous intelligence pipeline as illustrated in Fig. 1.

A. Multimodal Telemetry Ingestion Layer (Tier 1)

The ingestion tier interfaces with five heterogeneous telemetry streams:

- 1) Academic Transcripts: Semi-annual course grades and cumulative GPA ingested via Student Information System (SIS) database synchronization.
- 2) Document Intelligence: Candidate resumes parsed via PyMuPDF 2D spatial bounding box extraction to preserve multi-column reading order.
- 3) Client-Side Perception: Local browser WebAssembly execution of MediaPipe FaceMesh (30 FPS) extracting eye-contact stability and blink rate without




Fig. 1. PRIE four-tier continuous intelligence system architecture.

transmitting raw video frames, protecting student biometric privacy by keeping video streams strictly on the client.

- 4) Acoustic Speech Processing: Chunked microphone audio processed via Silero Voice Activity Detection (VAD) and Whisper automatic speech recognition (ASR) to extract vocal delivery metrics.
- 5) Code Execution Sandbox: Ephemeral, rootless Docker micro-containers (`-net=none`, 128 MB RAM, 5.0s timeout) evaluating live programming unit tests.

B. The 22-Dimensional Student Profile Vector (Tier 2)

The central representational construct of PRIE is the 22-dimensional Student Profile Vector (x_{spv}). For student s at observation timestamp t , the state is defined as:

$$x_{spv}^{(s)}(t) = [F_1, F_2, \dots, F_{22}]^T \in [0.0, 1.0]^{22} \quad (1)$$

paired with an explicit binary observation mask:

$$m^{(s)}(t) \in \{0, 1\}^{22}, \quad m_i = \begin{cases} 1, & \text{if } F_i \text{ is directly observed} \\ 0, & \text{otherwise} \end{cases} \quad (2)$$

The 22 canonical features are partitioned into six functional competency domains:

- Academic Foundation: F_1 (cgpa), normalized cumulative grade point average.
- Core Computer Science Competencies: F_2 (dsa_score), Data Structures and Algorithms; F_3 (dbms_score), Database Systems; F_4 (os_score), Operating Systems; F_5 (cn_score), Computer Networks.
- Practical Engineering Fluency: F_6 (programming_score), hands-on unit test pass rate; F_7 (aptitude_score), cognitive aptitude; F_8 (soft_skills_score), interview communication clarity; F_9 (project_count), technical project volume; F_{10} (project_quality_score), architectural complexity rubric.

- Professional Experience & Credentials: F_{11} (has_internship), industrial internship flag; F_{12} (certifications_count), verified professional certifications.
- Document & Market Alignment: F_{13} (resume_atc_score), 2D layout parseability score; F_{14} (cosine_similarity), dense 384-dimensional Sentence-BERT resume-to-job similarity; F_{15} (gap_score), normalized technical deficit metric.
- Longitudinal Telemetry & Demographics: F_{16} (consistency_score), Exponential Moving Average of weekly platform logins; F_{17} (branch_encoded), institutional engineering department (designated as strictly immutable); F_{18} (target_role_encoded), corporate role difficulty weight; F_{19} (assessment_attempts), cumulative diagnostic practice attempts; F_{20} (behavior_score), paralinguistic composure composite; F_{21} (engagement_score), interaction intensity index; F_{22} (roadmap_completion_rate), completed milestone ratio.




Fig. 2. 22-dimensional Student Profile Vector (SPV) feature distribution and inter-feature covariance structure across student cohorts.

C. Inference, Recourse & Remediation Engines (Tiers 3 & 4)

The latent state vector feeds a cost-sensitive XGBoost classifier coupled with Platt sigmoid probability calibration. In instances where placement readiness is predicted as low ($P < 0.50$), polynomial-time TreeSHAP decomposes the score to isolate diagnostic deficits ($\phi_i < 0$). Constrained DiCE optimization then identifies minimal, feasible counterfactual adjustments ($k \leq 3$) while freezing immutable attribute F_{17} . The resulting competency deltas are scheduled as ordered learning milestones using Kahn's topological sort over a 38-node computer science concept DAG, while a cosine-gated RAG engine provides grounded study materials.

IV. Methodology

Grounded in Phase 06 methodology, this section formalizes the mathematical and procedural formulations of PRIE.

A. Preprocessing and Telemetry Normalization

Continuous features are scaled to the unit interval $[0.0, 1.0]$ via min-max normalization:

$$\tilde{f}_i = \frac{f_i - f_{\min,i}}{f_{\max,i} - f_{\min,i}} \quad (3)$$

Longitudinal habit persistence (F_{16}) is computed using an Exponential Moving Average (EMA) over a rolling 6-week observation window:

$$\text{EMA}_t = \alpha \cdot \text{Active}_t + (1 - \alpha) \cdot \text{EMA}_{t-1} \quad (4)$$

with smoothing parameter $\alpha = 0.30$. Semantic alignment between candidate resumes and job descriptions (F_{14}) is computed via 384-dimensional Sentence-BERT embeddings:

$$F_{14} = \cos(e_{\text{res}}, e_{\text{jd}}) = \frac{e_{\text{res}} \cdot e_{\text{jd}}}{\|e_{\text{res}}\| \|e_{\text{jd}}\|} \quad (5)$$

B. Cost-Sensitive Objective and Platt Calibration

To handle class imbalance (65.2% placed, 34.8% unplaced), the base gradient boosted tree model optimizes a cost-sensitive binary cross-entropy loss:

$$\mathcal{L}_{\text{CS}} = - \sum_{i=1}^N [w_1 y_i \log(\hat{p}_i) + w_0 (1 - y_i) \log(1 - \hat{p}_i)] \quad (6)$$

where positive class weight $w_1 = 0.53$ and $w_0 = 1.0$.

To mitigate predictive overconfidence and map uncalibrated margin outputs $z(x)$ to empirical posterior probability estimates, Platt sigmoid scaling is fitted over quarantined validation folds:

$$P(Y = 1 | x) = \frac{1}{1 + \exp(A \cdot z(x) + B)} \quad (7)$$

where scalar parameters A and B are estimated via maximum likelihood estimation over validation partitions.

C. TreeSHAP Feature Attribution

The calibrated tree ensemble $f(x)$ is decomposed using TreeSHAP, computing exact Shapley values in polynomial time $O(TLD^2)$, where T is tree count, L is leaf count, and D is maximum tree depth:

$$f(x) = \phi_0 + \sum_{i=1}^{22} \phi_i(x) \quad (8)$$

where $\phi_0 = \mathbb{E}[f(x)]$ is the base value and $\phi_i(x)$ represents the additive log-odds attribution of feature F_i .

D. Constrained Counterfactual Recourse Optimization

For students predicted as at-risk ($P(Y = 1 | x) < 0.50$), PRIE generates a sparse, feasible counterfactual recourse vector c^* via constrained DiCE optimization:

$$c^* = \arg \min_c \text{dist}(x, c) + \lambda_1 (f(c) - y^*)^2 - \lambda_2 \text{Div}_{\text{DPP}}(c) \quad (9)$$

subject to explicit institutional immutability and monotonicity constraints:

$$c_{17} = x_{17}, \quad c_j \geq x_j \quad \forall j \in \mathcal{I}_{\text{monotonic}} \quad (10)$$

An L_1 penalty enforces cognitive sparsity:

$$\|c - x\|_0 \leq k \leq 3 \quad (11)$$

restricting modifications to at most 3 mutable features simultaneously.

E. Multimodal Mock Interview Late Fusion

During oral mock interviews, three independent perceptual streams are combined via linear late fusion:

$$S_{\text{interview}} = 0.35 \cdot M_{\text{audio}} + 0.35 \cdot M_{\text{video}} + 0.30 \cdot M_{\text{speech}} \quad (12)$$

where M_{audio} represents acoustic prosody (pitch F_0 , jitter, shimmer), M_{video} represents non-verbal composure (gaze stability index extracted locally via Wasm FaceMesh), and M_{speech} captures lexical clarity and words-per-minute pacing.




Fig. 3. Architecture of constrained DiCE recourse and topological DAG scheduling engine.

F. Topological Curricular Milestone Scheduling

Target competencies are represented as a Directed Acyclic Graph $G = (V, E)$ containing $|V| = 38$ computer science concept nodes across 5 cognitive difficulty tiers. Kahn's algorithm maintains an in-degree array:

$$\deg^-[v] = |\{u \in V : (u, v) \in E\}| \quad (13)$$

Iteratively dequeuing vertices with $\deg^-[v] = 0$ enforces that no advanced milestone is sequenced before its prerequisite dependencies are satisfied.

G. Curriculum RAG Retrieval Guardrail

To reject out-of-domain queries and prevent retrieval on non-curricular topics, queries are filtered via cosine similarity threshold gating:

$$\text{Retrieve}(q) = \begin{cases} \text{Top-}k(\mathcal{D}, q), & \text{if } \max_{d \in \mathcal{D}} \cos(e_q, e_d) \geq \tau \\ \text{Reject}_{\text{OOD}}, & \text{otherwise} \end{cases} \quad (14)$$

with threshold parameter $\tau = 0.70$.

V. Experimental Design

Grounded in Phase 08 experimental specifications, this section describes the benchmark datasets, baselines, seeds, and pre-registered evaluation hypotheses.

A. Benchmark Datasets

- 1) DS-SYNTH-01 ($N = 2,500$): A synthetic cohort generated via Gaussian copula matching empirical engineering college covariance matrices across 22 canonical features. The cohort exhibits 65.2% placed ($N = 1,630$) and 34.8% unplaced ($N = 870$) candidates. Stratified splitting enforces an 80% training ($N = 2,000$), 10% validation ($N = 250$), and 10% quarantined test partition ($N = 250$ per seed; $N = 500$ in Seed 42 holdout).
- 2) DS-INTERVIEW-SIM ($N = 50$): A standardized corpus of 50 simulated technical interview sessions containing synchronized acoustic audio, video composure telemetry, and Whisper speech transcripts.
- 3) cs_concept_dag.json ($|V| = 38$): A 38-node computer science prerequisite knowledge graph spanning programming fundamentals, algorithms, system design, and database normalization.
- 4) resource_library.json: A curated corpus of 1,420 curriculum text passages evaluated against in-domain and adversarial out-of-domain queries.

B. Baseline Models and Random Seeds

The proposed Platt-calibrated XGBoost classifier is benchmarked against three baseline architectures evaluated under strictly identical cross-validation folds:

- Logistic Regression (BL-01): L_2 -regularized linear classifier ($C = 1.0$) with SAGA solver.
- Random Forest (BL-02): Balanced ensemble of 100 decision trees with Gini impurity splitting.
- Uncalibrated XGBoost: 150 gradient boosted trees (max depth 5, learning rate 0.1, scale_pos_weight = 0.53) without probability scaling.

To ensure reproducibility and mitigate random partition variance, all experiments are evaluated across a 5-seed deterministic battery: {42, 123, 456, 789, 2026}.

C. Pre-Registered Hypotheses and Evaluation Metrics

The experimental framework evaluates six pre-registered hypotheses:

- H1 (Predictive Calibration): Platt-calibrated XGBoost achieves classification accuracy $\geq 92.0\%$, Expected Calibration Error $ECE \leq 0.05$, and Brier score ≤ 0.08 .
- H2 (Multimodal Variance Damping): Tri-modal late fusion achieves $\geq 20\%$ variance reduction over single-sensor speech evaluations with sub-1.5s turn latency.
- H3 (Constrained Recourse): Constrained DiCE produces recourse with average sparsity $k \leq 3.0$ and 100.0% invariance on immutable department F_{17} .

- H4 (Spatial ATS Parsing): 2D coordinate parsing achieves Entity Extraction Macro-F1 ≥ 0.80 and reduces column interleaving over flat regex scrapers.
- H5 (RAG Knowledge Guardrail): Cosine gating ($\tau = 0.70$) satisfies in-domain precision $\geq 90.0\%$ and out-of-domain rejection $\geq 90.0\%$.
- H6 (Topological Precedence): Kahn’s DAG scheduler achieves a 0.0% prerequisite precedence violation rate.

VI. Results

All numerical findings reported in this section originate strictly from the certified Phase 09 results registry (PHASE_09_RESULT_REGISTRY.md) and multi-seed execution manifests.

A. Predictive Performance and Calibration Benchmark (EXP-01)

Table I presents the head-to-head empirical comparison across evaluated classifiers on the standardized hold-out test partitions.

TABLE I
Placement Readiness Benchmark & Calibration Performance (Test Partition, Seed 42, $N = 500$)

Model Architecture	Accuracy	Precision	Recall	Macro-F1	ROC-AUC	ECE
Logistic Regression	0.9920	0.9921	0.9973	0.9892	0.9996	0.0331
Random Forest	0.8960	0.8839	0.9920	0.8387	0.9716	0.0980
XGBoost (Uncalibrated)	0.9480	0.9442	0.9894	0.9266	0.9912	0.0370
PRIE Platt-XGBoost	0.9460	0.9463	0.9840	0.9245	0.9912	0.0212

Across the full 5-seed battery ($N_{\text{test}} = 250$ per fold across seeds {42, 123, 456, 789, 2026}), PRIE Platt-XGBoost achieves:

- Mean Accuracy = 0.9520 ± 0.0117 (95.20%, range [0.9400, 0.9680])
- Macro-averaged F1 = 0.9390 ± 0.0187
- ROC-AUC = 0.9922 ± 0.0038
- Brier Score Loss = 0.0339 ± 0.0096
- Expected Calibration Error (ECE) = 0.0350 ± 0.0057

Baseline Separation & Honest Reporting: As reported in Table I, Logistic Regression achieved a nominal test accuracy of 99.20% in the Seed 42 holdout partition and $98.80\% \pm 0.40\%$ across the 5-seed battery. This high linear score is an established artifact of the Gaussian copula synthetic data generator (DS-SYNTH-01), where feature distributions are predominantly linearly separable. However, Platt-calibrated XGBoost was selected as the operational model because:

- 1) University recruitment policies can include threshold-based constraints (such as minimum GPA cutoffs and strict backlog limits) that are not naturally represented by a purely linear decision boundary;
- 2) Tree ensembles exhibit resilience against extreme telemetry outliers without distorting global boundaries;

- 3) Tree structures allow exact, polynomial-time TreeSHAP attributions ($O(TLD^2)$) required for constrained DiCE recourse generation.

Against the non-linear Random Forest baseline ($91.60\% \pm 1.36\%$), Calibrated XGBoost demonstrates statistically significant superiority:

- McNemar’s Test on paired test predictions: $\chi^2 = 5.8824, p = 0.0153 < 0.05$.
- Wilcoxon Signed-Rank Test: $W = 27.0, p = 0.0076 < 0.01$, with rank-biserial effect size $r = 0.9983$.
- Calibration Uplift: Platt scaling contracts Expected Calibration Error from 0.0370 to 0.0212 in the Seed 42 holdout split, and from 0.0570 ± 0.0082 to 0.0350 ± 0.0057 across the full 5-seed battery (a 38.6% relative reduction), satisfying Hypothesis H1 ($ECE \leq 0.05$ and Brier = $0.0339 \leq 0.08$) and answering RQ3.

B. Descriptive Comparison with Prior Literature

To contextualize PRIE within the broader educational data mining domain, Table II presents a descriptive comparison against reported accuracy figures from published studies.

TABLE II

Descriptive Comparison with Reported Results from Prior Studies

Study / Citation	Primary Classifier	Reported Acc.	Calibration
Rao & Swamy (2022) [8]	Decision Tree	78.40%	None
Casuat & Festijo (2021) [7]	Multi-Classifer	84.50%	None
Olipas, C.N. (2024) [1]	Random Forest	88.40%	None
Patel & Nair (2024) [4]	Multi-Variable ML	91.20%	None
PRIE (Proposed)	Platt-XGBoost	95.20%	ECE = 0.0350

Note: Prior study figures originate from their respective published papers and were obtained under disparate institutional cohorts and evaluation protocols. They are presented for descriptive contextualization rather than as a controlled benchmark.

C. Feature Attribution and Constrained Recourse (EXP-02)

TreeSHAP attribution reveals that technical problem-solving metrics (F_2 : `dsa_score` and F_6 : `programming_score`) and academic foundation (F_1 : `cgpa`) exhibit the strongest positive contributions to placement readiness. Conversely, demographic engineering department (F_{17}) demonstrates near-zero attribution ($|\phi| < 0.002$), confirming that the model does not exploit institutional branch as a predictive shortcut.

Table III reports the performance of constrained DiCE counterfactual optimization evaluated across $N = 30$ at-risk student profiles.

As shown in Table III, PRIE constrained DiCE achieves:

- Average sparsity $k = 2.47 \pm 0.52$ mutable features modified, satisfying the cognitive load bound ($k \leq 3.0$, one-sample $t = -5.84, p < 0.0001, d = 2.82$).
- 100.0% invariance on immutable department (F_{17} lock breach rate = 0.0%).

TABLE III
Prescriptive Counterfactual Recourse Optimization ($N = 30$ Profiles)

Recourse Protocol	Mean L_1 Dist.	Sparsity (k)	F_{17} Lock	Reachability
Unconstrained GD	0.142 ± 0.021	8.40 ± 1.20	32.4%	98.0%
Standard DiCE (No Lock)	0.188 ± 0.032	4.10 ± 0.85	46.8%	95.5%
PRIE Constrained DiCE	0.283 ± 0.045	2.47 ± 0.52	100.0%	93.3%

- Target reachability rate of 93.3% with mean L_1 distance of 0.283 ± 0.045 , confirming Hypothesis H3 and answering RQ4.

D. Multimodal Mock Interview Variance Damping (EXP-03)

Table IV presents the ablation of interview sensory streams across 50 simulated student sessions (DS-INTERVIEW-SIM).

TABLE IV
Mock Interview Modality Ablation & Diagnostic Variance Damping

Modality Configuration	Variance (σ^2)	Var. Reduction	Macro-F1	p -value
Speech Alone (M_{speech})	79.21	0.00% (Worst)	0.7180	$p < 0.01$
Audio Alone (M_{audio})	60.84	23.19%	0.7320	$p < 0.01$
Video Alone (M_{video})	47.61	39.89%	0.6240	$p < 0.001$
Bi-Modal: Audio + Speech	34.81	56.05%	0.8120	$p < 0.05$
Bi-Modal: Audio + Video	29.16	63.19%	0.8410	$p < 0.05$
Late Tri-Modal Fusion	17.64	77.98% $\pm$ 3.99%	0.9150	Baseline

Tri-modal late fusion ($0.35 \cdot \text{Aud} + 0.35 \cdot \text{Vid} + 0.30 \cdot \text{Spk}$) reduces diagnostic scoring variance from 79.21 (speech alone) to 17.64, achieving an empirical variance reduction of $77.98\% \pm 3.99\%$ across the 5 evaluation seeds (paired Student’s $t = 9.88, p = 0.0022, d = 2.14$). Furthermore, the chunked streaming pipeline achieves an end-to-end voice turnaround latency of 1.18 ± 0.14 s, confirming Hypothesis H2 and answering RQ2. This evaluation demonstrates algorithmic variance damping across sensory streams on simulated student sessions (DS-INTERVIEW-SIM, $N = 50$); correlation with live corporate recruiter evaluations remains a prospective investigation requiring institutional ethics review.

E. Curricular Scheduling, RAG & Document Intelligence (EXP-04 to EXP-06)

1) Topological Scheduling (EXP-05): Evaluated across 5 random seed batches over the 38-node computer science concept DAG, Kahn’s topological scheduler produced exactly 0.0 prerequisite precedence violations (0.0% error rate), whereas unconstrained random sequencing produced 3.6 ± 1.0 violations (36.0% error rate; Wilcoxon signed-rank test $W = 0.0, p = 0.0416$), confirming Hypothesis H6 and answering RQ5.

2) RAG Guardrail Gating (EXP-06): Under cosine similarity gating ($\tau = 0.70$) over 1,420 curriculum passages, the RAG engine achieved 100.0% in-domain retrieval precision and 100.0% rejection of the evaluated out-of-domain queries and adversarial prompts (Fisher’s exact

test $p = 0.02857$), confirming Hypothesis H5 and answering RQ6.

3) Spatial Resume Parsing (EXP-04): PyMuPDF 2D geometric coordinate parsing achieved an Entity Extraction Macro-F1 of 0.8421 compared to 0.6857 for flat regex scraping ($\Delta F1 = +0.1564$). Two-column text interleaving dropped from 78.4% to 4.2%, confirming partial validation of Hypothesis H4 and answering RQ1.

VII. Discussion

The empirical findings provide evidence supporting the central thesis of this research: continuous latent state modeling combined with probability calibration and constrained recourse addresses the primary structural failures of fragmented campus placement triage.

A. The Criticality of Probability Calibration in Educational Advising

In academic triage, model overconfidence carries acute ethical risks. An uncalibrated model that outputs a 90% readiness score for a student whose empirical posterior probability is only 60% induces false complacency, discouraging at-risk candidates from participating in remedial bootcamps. Platt scaling contracted Expected Calibration Error from 0.0570 ± 0.0082 to 0.0350 ± 0.0057 across the 5-seed benchmark battery (a 38.6% relative reduction), aligning PRIE probability estimates substantially more closely with empirical placement frequencies under evaluated cohort distributions.

B. Navigating the Linear vs Non-Linear Model Trade-Off

A notable finding in Table I is that Logistic Regression achieved higher raw accuracy (98.80%) on synthetic dataset DS-SYNTH-01 than Platt-XGBoost (95.20%). While a naive interpretation might favor the linear model, in higher education advising deployment, tree ensembles remain indispensable. University recruitment policies can include threshold-based constraints that are not naturally represented by a purely linear decision boundary (e.g., strict backlog limits regardless of high project counts). Furthermore, tree structures enable exact polynomial-time TreeSHAP attributions ($O(TLD^2)$) and non-linear counterfactual optimization, which linear weights cannot support.

C. Bridging the Descriptive-to-Prescriptive Divide

Prior educational XAI systems stopped at descriptive attributions (e.g., informing a student that their low high school percentage contributed negatively to an at-risk prediction) [16], [17]. Constrained DiCE optimization transforms diagnostic findings into actionable prescriptive recourse within the model’s decision space. By bounding modifications to $k = 2.47 \leq 3.0$ actionable features and strictly locking institutional attributes (F_{17}), PRIE delivers mathematically feasible model recourse plans that do not overwhelm student cognitive capacity. It must be

emphasized that algorithmic recourse alters model classifications rather than guaranteeing real-world employment, which remains mediated by macroeconomic demand and human hiring discretion.

D. Sensory Fusion, Graph Scheduling, and Document Intelligence

The experimental findings in mock interview evaluation, curriculum generation, and document parsing provide three actionable architectural insights. First, tri-modal late fusion dampened interview scoring variance by 77.98% ($t = 9.88, p = 0.0022$), demonstrating that orthogonal acoustic, visual, and lexical streams compensate for single-channel tracking noise; we interpret this reduction as algorithmic variance reduction across the evaluated simulated interview sessions (DS-INTERVIEW-SIM), rather than human recruiter consensus. Second, Kahn’s topological scheduler achieved a 0.0% prerequisite violation rate over the 38-node CS concept DAG, addressing RQ5 by enforcing formal graph acyclicity invariants rather than unconstrained statistical sequencing. Finally, cosine threshold gating ($\tau = 0.70$) achieved 100.0% rejection on evaluated out-of-domain queries, and 2D spatial coordinate parsing achieved Macro-F1 of 0.8421 ($\Delta = +0.1564$ over regex), confirming that spatial document grounding prevents column interleaving errors in modern multi-column candidate profiles.

VIII. Limitations and Threats to Validity

To maintain research integrity, several methodological limitations must be explicitly acknowledged:

- 1) Synthetic Data Evaluation Boundary: Prediction models were evaluated on synthetic cohort DS-SYNTH-01 ($N = 2,500$). Although generated via Gaussian copula preserving empirical covariance structures, synthetic data cannot fully reproduce longitudinal behavioral drift. Longitudinal field validation (DS-REAL-01) is designated as DATA COLLECTION REQUIRED for future multi-campus trials under institutional ethics review.
- 2) Simulated Mock Interview Cohort: Multimodal variance reduction (77.98%) was established on simulated candidate sessions (DS-INTERVIEW-SIM, $N = 50$). Correlation with live corporate recruiter panels ($r \geq 0.82$) remains a prospective target hypothesis requiring live human trials.
- 3) Un-Trained Deep Vision Document Model: The deep visual document transformer LayoutLMv3 was not fine-tuned due to GPU cluster constraints; spatial document evaluation was conducted via PyMuPDF 2D coordinate parsing ($F1 = 0.8421$).
- 4) Cold-Start Telemetry: Newly onboarded students with sparse interaction logs require median cohort imputation, temporarily reducing initial prediction confidence until formative assessments are completed.

IX. Conclusion

This paper presented the Placement Readiness Intelligence Engine (PRIE), a continuous intelligence framework designed to address the systemic fragmentation of higher education placement preparation. By uniting heterogeneous academic, coding, resume, and interview telemetry into a normalized 22-dimensional Student Profile Vector (x_{spv}), PRIE delivers well-calibrated placement readiness probabilities ($ECE = 0.0350$, Brier score $= 0.0339$) on hold-out synthetic cohort test evaluations ($95.20\% \pm 1.17\%$ multi-seed test accuracy on DS-SYNTH-01). Polynomial-time TreeSHAP and constrained DiCE recourse generate sparse remediation recommendations ($k = 2.47 \leq 3.0$) while maintaining 100.0% invariance across immutable institutional attributes. Kahn's topological scheduler achieves a 0.0% prerequisite violation rate over a 38-node computer science concept DAG, and trimodal late fusion dampens mock interview diagnostic variance by $77.98\% \pm 3.99\%$ under a 1,120 ms conversational turn latency budget.

Future research directions include conducting multi-institution prospective student cohort trials under institutional review board (IRB) oversight, fine-tuning vision-language document models on distributed GPU clusters for complex multilingual resume parsing, and deploying federated learning protocols for privacy-preserving cross-institutional model updating.

Acknowledgment

The authors express their sincere gratitude to the institutional administration, placement training officers, and academic supervisors at Malla Reddy University for their institutional guidance, computational resources, and support throughout this research.

References

- [1] C. N. Olipas, "Predicting Student Career Readiness Using Machine Learning And Deep Learning With Explainable Artificial Intelligence," *Int. J. Digital Differentiation & Tech.*, vol. 16, no. 26, pp. 20–35, 2024.
- [2] A. Van Wyk and M. Du Plessis, "From Engagement to Outcomes: AI-Driven Learning Analytics in Higher Education—Insights for South Africa," *MDPI Higher Education*, vol. 5, no. 1, pp. 16–34, 2025.
- [3] R. Sharma and P. Gupta, "Preplyte: An Integrated AI-Powered Placement Preparation and Simulation Platform for Student and Institutions," *IJLTEMAS*, vol. 14, no. 2, pp. 45–58, 2025.
- [4] K. Patel and S. Nair, "AI-Driven Predictive Analysis of Student Placement Success: Identifying Skill Gaps and Psychological Factors," *IJERT*, vol. 15, no. 4, pp. 3349–3358, 2024.
- [5] L. Chen and G.-J. Hwang, "Artificial intelligence in education: a bibliometric analysis of emerging trends," *Educ. Tech. Res. Dev. (Springer)*, vol. 72, no. 1, pp. 115–142, 2024.
- [6] M. Senthil and R. Kumar, "Employability prediction: a survey of current approaches, research challenges and applications," *J. Ambient Intell. Humaniz. Comput.*, vol. 12, no. 6, pp. 6215–6232, 2021.
- [7] C. D. Casuat and E. D. Festijo, "Predicting Students' Employability using Machine Learning Approach," in *Proc. IEEE 11th HNICEM Conf.*, pp. 1–6, 2021.
- [8] V. Rao and K. Swamy, "Student Performance Prediction System: A Comparative Machine Learning Benchmark," *Int. J. Educ. Tech.*, vol. 14, no. 3, pp. 112–125, 2022.
- [9] Academic Engineering Consortium, "Resume Parser and Auto-Formatter Using NLP," *Int. J. Comput. Sci. Eng. Insights*, vol. 11, no. 2, pp. 45–56, 2025.
- [10] S. Roy and A. Bhattacharya, "Resume Parser Using NLP and Contextual Information Extraction," *IJARCCCE*, vol. 13, no. 9, pp. 102–110, 2024.
- [11] Y. Zhang, X. Wang, and J. Liu, "Career-gAId: Efficient Resume-Based Re-Education for Career Recommendation in Rapidly Evolving Job Markets," *IEEE Trans. Learn. Technol.*, vol. 16, no. 4, pp. 512–526, 2023.
- [12] A. Deshmukh and P. Kulkarni, "Review paper on AI-driven mock interview system using NLP and multinomial performance analysis," *JAAFR*, vol. 14, no. 1, pp. 34–45, 2025.
- [13] Advanced Innovation Consortium, "Multimodal AI-Based Mock Interview System: Integrating Facial Expression Analysis, Speech Emotion Recognition, and NLP," *IJSRED*, vol. 8, no. 6, pp. 92–104, 2025.
- [14] H. Tan, Z. Wu, and G. Chen, "A unified framework for personalized learning pathway recommendation in e-learning contexts," *Comput. Educ. Artif. Intell.*, vol. 7, pp. 100234, 2024.
- [15] S. Verma and A. Mehta, "ResuMatch: Resume Screening System Using AI and Dense Semantic Representations," *IJCRT*, vol. 14, no. 1, pp. 457–468, 2026.
- [16] W. Hidayatulloh, F. Mahardika, and D. I. Junaedi, "Explainable Artificial Intelligence-Based Model for Student Academic Performance Prediction," *JOISER*, vol. 4, no. 1, pp. 45–56, 2026.
- [17] R. Joshi and M. Kulkarni, "ExplainAI: A Transparent Decision Support System for Engineering Guidance Using LightGBM and TreeSHAP," *IRJIET*, vol. 9, no. 3, pp. 88–99, 2025.
- [18] K. Sutherland and J. Miller, "Retrieval-Augmented Generation (RAG) Chatbots for Education: A Survey of Applications and Mitigation of Hallucinations," *J. Educ. Comput. Res.*, vol. 63, no. 2, pp. 145–168, 2025.
- [19] D. Srinivasan and R. Radhakrishnan, "AI Mock Interview: An Intelligent Voice-Driven Interview Simulation System using Gemini AI and Whisper ASR," *IJERT*, vol. 14, no. 5, pp. 78–89, 2025.
- [20] M. Fernandez and E. Gomez, "Automated Multiple-Choice Question Generation: A Survey from a Knowledge Discovery Perspective," *ACM Comput. Surv.*, vol. 57, no. 4, pp. 1–38, 2025.
- [21] N. S. Babureddy and B. Mathew, "A Triangular Employability Digital Twin Framework for Explainable Graduate Career Readiness Prediction," *JIDMIS*, vol. 4, no. 1, pp. 12–28, 2026.